



Chapter 1 — Substrate

What a language model is, and what one run of it computes

The answer arrives a piece at a time

animation

a question typed, and the reply appearing left to right at natural speed

1600 x 422 px

The text you send is the **prompt**. It writes left to right, in pieces — and if you ask the same question twice, the two answers differ.

observed

Substrate

Formation

Control

Measure


Failure

Limits

Language models are one branch of AI

Every band contains the one inside it. Not all artificial intelligence is a language model.

dashed = a boundary drawn by usage, not by definition — "artificial intelligence" is applied inconsistently



Large language model — a neural network fitted to predict text. Shortened to **LLM**.

definition

Substrate

Formation

Control

Measure


Failure

Limits

Reliability follows how much was written

Reliability tracks how much has been written, not how important the topic is.

roughly where your own research
literature sits



The corpus decides where it is strong. Nobody has trained a model mainly on your literature.

definition



A model is numbers plus a program

the parameters

one long list of numbers

-0.0412	0.9037	0.1188	-1.2740
0.0006	0.4451	-0.3319	0.7702
0.0925	-0.0187	1.1046	-0.6538

... and so on, counted in billions

the program

the arithmetic that uses them

1. read the numbers from the list
2. multiply and add them against the input, layer after layer
3. hand back one score per token

Both are files on a disk. Together they are the model.

Parameters

also called **weights**

Network

the layered arithmetic

Size

how many parameters

definition

Substrate

Formation

Control


Measure

Failure

Limits

Training set those numbers and then stopped

Nothing to the right of the line changes anything to the left of it.



Training

the numbers are being changed

Inference

running the finished numbers on
your text

Knowledge cutoff



the date training stopped

definition



Training searches for the smallest total error

$y = 2x + 1$ exactly — chosen so slope 2.00 and intercept 1.00 can be verified by hand. A real fit does not land on zero residual.



Two parameters here: **slope 2.00** and **intercept 1.00**. A language model does the same search over billions of them, and the data is text.



computed

Training minimises surprise at the next token

1. a training example

The beam failed in shear because
the stirrups were L

during training, the true
next token is known



You cannot define "surprised" without a probability. The distribution is what training shaped.

You cannot define *surprised* without a **probability**. So the thing training shaped is a distribution over every **token** — a chunk of text from a fixed list, defined on slide 12 — that could come next.

derived

Substrate

Formation

Control


Measure

Failure

Limits

Training changed the numbers, nothing else

ILLUSTRATIVE values, drawn from a fixed seed. The movement is the point, not the numbers.



Neural network

layers of multiply-and-add,
named after a biological analogy

Learning

the search you just watched


The brain comparison

correlational; the appendix
carries what it establishes

[V] Schrimpf, M., Blank, I. A., Tuckute, G., et al., 2021. *PNAS* 118(45):e2105646118, DOI 10.1073/pnas.2105646118.

[V] Hadidi, N., Fegghi, E., Song, B. H., Blank, I. A. & Kao, J. C., 2026. *Nature Communications* 17(1):5769, DOI 10.1038/s41467-026-72253-7.


Text becomes numbers, then scores, then text



Every box gets its own slide, in this order. **Token** — a chunk of characters from a fixed list, defined next.

definition

Everything given sits in the context window



Re-read in full every time the network runs. **Standing instructions** — text placed in the window at the start of every session.

definition

Substrate

Formation

Control


Measure

Failure

Limits

Tokens come from a fixed list

SCHEMATIC — no tokeniser was run. Capture a real split before delivery.



marks a leading space, which belongs to the token. "shear" and "stirrups" split mid-word — eleven token IDs, in slate, for nine words.

* Not only text. An image or a PDF page is cut into small squares and turned into numbers in the same window.

The vocabulary

tens of thousands of entries

Token IDs

labels, not quantities


No letters

the network is handed integers

schematic



Token count follows frequency, not length



Billed
per token

Context limit
counted in tokens


Speed
quoted per token

schematic



The network returns one score per token

illustrative logits — the softmax transform applied to them is exact, over these eight plus one term for the rest of the vocabulary



The raw scores are **logits**. One run of the network over the window is one **pass**. The four shown do not add to one — the rest of the vocabulary carries the remainder.

definition

Substrate

Formation


Control

Measure

Failure

Limits

Attention decides which earlier tokens matter



animation A3 replaces this figure — 1600 × 600 px

A **weighted sum**: the weight on each earlier position comes from comparing it against the position being scored. **Transformer** — the arrangement of arithmetic built from it.

[V] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L. & Polosukhin, I., 2017. *Attention Is All You Need*, arXiv:1706.03762.

Softmax turns scores into probabilities

$$p_i = \frac{\exp\left(\frac{\text{score for token } i}{\text{temperature}}\right)}{\underbrace{\sum_j \exp(z_j/T)}_{\text{sum over every token, so the result adds to 1}}}$$

Worked case — restricted to two candidates so the sum can be checked by hand.

$$z_{\text{absent}} = 2.0 \quad z_{\text{corroded}} = 1.0$$

$$p = \frac{e^{2.0}}{e^{2.0} + e^{1.0}} = \frac{7.389}{10.107} = 0.731$$

All quantities dimensionless.

computed

Always taking the top token repeats

animation

temperature zero, run long enough that the output falls into a loop

1600 x 422 px

Same input, same scores, same top token — so a recurring phrase recreates the context that produced it. Taking the top every time is **greedy** selection.

derived

Substrate

Formation

Control


Measure

Failure

Limits

The token is drawn, not taken

illustrative logits — the softmax transform applied to them is exact, over these eight plus one term for the rest of the vocabulary



animation A4 replaces this figure — 1600×600 px — the bars must be identical on every draw

The bars are computed from the frozen numbers and never change. The **draw** does.

derived

Substrate

Formation

Control


Measure

Failure

Limits

Temperature sets how sharply the top score wins

Same eight candidates, four temperatures — logits are $\log(\text{table probability})$, so softmax at $T = 1$ reproduces the table exactly



animation A5 replaces this figure — 1600×450 px

Low T

one token takes almost
everything

$T = 1$

the scores used as they are

High T

the **probability** spreads

computed

Substrate

Formation

Control


Measure

Failure

Limits

Each token is appended and everything reruns

ILLUSTRATIVE — the same frozen numbers are read on every pass. Only the text handed to them has grown.




animation A7 replaces this figure — 1600×450 px

Four passes here. A paragraph is a few hundred. The **parameters** are unchanged on every one of them.

schematic




Two runs give different answers



The puzzle from the first slide, resolved. Everything left of **draw one** is fixed.

derived


Tools write into the window



A **tool** is a program the system runs alongside the network. Its result is written into the window as tokens, and the pass that follows is the one already described.

[V] Anthropic, 2026. *Code execution tool*, platform.claude.com — vendor documentation, fetched 2026-08-20.

The session ends and the window goes



Discarded

the window, the conversation, anything you
pasted

Unchanged

the **parameters**, in every session

definition



What we covered, and what is still open

What we covered

- ▶ A model is a list of **numbers** and a program that reads them
- ▶ Training fixed those numbers and stopped
- ▶ What it minimised was surprise, so the output is a **distribution**
- ▶ Text is cut into **tokens** from a fixed list
- ▶ Every pass scores every token, then draws one
- ▶ It sees only the context window, and discards it

What is still open

- ▶ Nothing in scoring text makes a thing that answers a question, follows an instruction, or declines a request. What was added?
- ▶ Why does it agree with you so readily?
- ▶ If it is reliable where much has been written, how do you tell which regime you are in?
- ▶ How would you ever know a change to your prompt actually helped?

derived