
[--] *No claim is made on this slide.*

A model that only scores text continues your question instead of answering it

animation


the running question typed into a model that has had
nothing added — it continues with more questions

1600 x 422 px

The question, throughout this chapter: ’

[--] *Direct observation, made by running it. No external source.*


Four stages turn a predictor into something that answers



Every stage after the first is a person's judgement, written down and turned into a number. This chapter walks the diagram once, left to right, and stops at each stage to ask what the job actually looks like.

[--] *Follows from earlier slides in this chapter. No external source.*

Pretraining runs one objective over a very large corpus, and then it stops



No labels

the text supervises itself, which
is why the corpus could be so
large

Knowledge cutoff


the date it stopped, and a
consequence rather than a policy

The corpus decides

where it is strong, and where it
is not

[--] Follows from earlier slides in this chapter. No external source.

People write the answer that should have come out, and the model is trained on the pair



Supervised fine-tuning

the model learns the *shape* of an answer — that a question gets answered

Cheap in data terms

and it does most of the visible work of turning autocomplete into an assistant


[V] Ouyang, L., Wu, J., Jiang, X., et al., 2022. *NeurIPS 2022*, arXiv:2203.02155.

The person repairs the text as well as judging it, against a written specification

The specification is not about being right. It is about punctuation, hedging, person, sentence length and the shape of a list.

[- -] *Structural observation from five annotation instruments held by the instructor. Not citable, and never becomes a citation. Three of five documents.*


In two of the five instruments, the repaired text is what gets collected



The labour recurs in three documents. The repaired text entering the data is a separate claim on weaker evidence, and in one instrument the repaired version is explicitly not trained on.

[- -] Structural observation from five annotation instruments held by the instructor. Not citable, and never becomes a citation. Two of five documents for collection; three of five for the labour.

Demonstrations cannot cover everything, so preferences are collected instead



Writing the ideal answer to every possible request is impossible. Judging which of two answers is better is not.

[V] Stiennon, N., Ouyang, L., Wu, J., et al., 2020. *preprint — venue unconfirmed*, arXiv:2009.01325.